A Constraint Opinion Model *

Fabio Gadducci[1], Carlos Olarte[2], and Frank Valencia[3,4]

[1] University of Pisa, Italy `fabio.gadducci@unipi.it`
[2] LIPN, CNRS UMR 7030, Université Sorbonne Paris Nord, France
`olarte@lipn.univ-paris13.fr`
[3] CNRS-LIX, École Polytechnique de Paris, France
[4] Department of Electronics and Computer Science, Pontificia Universidad Javeriana, Cali,
Colombia `frank.valencia@gmail.com`

Abstract. This paper introduces a generalized opinion model that extends the standard DeGroot model by representing agents' opinions and influences as soft constraints rather than single real values. This allows for modeling scenarios beyond the scope of the DeGroot model, such as agents sharing partial information and preferences, engaging in discussions on multiple topics simultaneously, and representing opinions with different degrees of uncertainty. By considering soft constraints as influences, the proposed model captures also situations where agents impose conditions on how others' opinions are integrated during belief revision. Finally, the flexibility offered by soft constraints allows us to introduce a novel polarization measure that takes advantage of this generalized framework.

Keywords: Opinion Models · Soft constraints · Multi-Agent Systems · Social Networks · Cognitive Bias · Consensus

1 Introduction

Social networks play a significant role in *opinion formation*, consensus building and polarization among their users. The dynamics of opinion formation in such networks typically involve individuals sharing their views with their contacts, encountering different perspectives, and adjusting their beliefs in response. Models of opinion formation (see e.g. [11,12,1,3]) capture these dynamics to simulate and reason about opinion evolution.

The DeGroot model [11] is one of the most representative formalisms for opinion formation and consensus building in social networks. In this model, a social network is represented as a directed *influence graph*, where edges denote the weight, expressed as a real number, that an *agent* (i.e., an individual) carries on another. Each agent holds an opinion, also expressed as a real number, indicating their level of agreement with an underlying proposition. Agents repeatedly update their opinions by taking the weighted average of the opinions of those who influence them (i.e., their *neighbors* or *contacts*).

There is empirical evidence validating opinion formation through averaging in controlled sociological experiments (e.g., [8]). The DeGroot model is widely recognized for its elegant characterization of opinion consensus based on the topology of the influence graph and it remains a central focus of research for developing frameworks to understand opinion formation dynamics in social networks (e.g., [16,10,27,12,9,22,1,3]).

Nevertheless, when modeling real scenarios from social networks, we often have only *partial information* about agents' opinions and the influence they have on one another. In practice, opinion information may be incomplete or imprecise due to privacy constraints, self-censorship, or the dynamic nature of beliefs. Similarly, influence relationships are often difficult to quantify, as they depend on factors such as trust, authority, and exposure to diverse viewpoints, which are not always explicitly observable. This *uncertainty* complicates the application of classical models like DeGroot, which assume for all agents, fully known opinions and influences, typically represented as real values.

In this paper we introduce *Constraint Opinion Models*, a framework where both opinions and influences are represented as *(soft) constraints* rather than exact real values. This allows us to reason about meaningful situations where only partial information or preferences are available. We thus generalize the DeGroot model, at the level of opinions and influences, while keeping much of its mathematical simplicity.

The use of soft constraints to represent opinions offers several advantages. First, they seamlessly represent opinions on different topics or propositions, enabling the analysis of network behavior as agents discuss various subjects. Second, they allow for the representation of uncertainty and partial information. The is particularly important to model situation where an agent's exact opinion is unknown. Third, they support some forms of epistemic modeling, capturing beliefs where agents hold opinions about other agents. Additionally, the framework can express complex opinions that a single value cannot adequately represent, such as "extreme" viewpoints (e.g., agents that prefer any extreme option than a moderate one).

Regarding the representation of influences, soft constraints provide an extra flexibility. They enable the definition of "filters" that impose boundaries on how agents adjust their opinions while preserving their core beliefs. Moreover, soft constraints allows for representing conditional influences, where the weight of the influence depends on the incoming information or the subject being discussed.

Finally, a key challenge in social system analysis is measuring the difference between two opinions as this is the basis of any *polarization* measure [13]. In the DeGroot model, where opinions are real values, this is straightforward. To extend this capability to constraint-based opinion models, we introduce a notion of distance between constraints. This allows for quantifying opinion divergence and assessing how polarized a system of agents sharing constraints becomes.

Organization. After recalling the basic notions of opinion models and semiring-based constraints in Section 2, we define constraint opinion models in Section 3. Section 4 is devoted to illustrate with several examples the possibilities offered by our framework in modeling scenarios beyond the classical DeGroot opinion model. Section 6 concludes the paper. The experiments shown in Section 4 can be reproduced with the aid of a tool we available at `https://github.com/promueva/constraint-opinion-model`.

2   Preliminaries

This section recalls some results on semirings, which are the algebraic structures adopted here for modeling (soft) constraints (Section 2.1). We also recall the notion of opinion models and belief revision in the standard DeGroot model (Section 2.2)

2.1   Monoids, Semirings and Soft Constraints

We shall use the values of a monoid, that can be combined, to represent *preferences*.

Definition 1 (Monoids, groups). *A (commutative) monoid is a triple $\langle A, \oplus, \mathbf{0} \rangle$ such that $\oplus : A \times A \to A$ is a commutative and associative function and $\mathbf{0} \in A$ is its* identity *element, i.e. $\forall a \in A.a \oplus \mathbf{0} = a$. A group is a four-tuple $\langle A, \oplus, \ominus, \mathbf{0} \rangle$ such that $\langle A, \oplus, \mathbf{0} \rangle$ is a monoid and $\ominus : A \times A \to A$ a function satisfying $\forall a \in A.a \ominus a = \mathbf{0}$.*

As usual, we use the infix notation: $a \otimes b$ stands for $\otimes(a, b)$.

Definition 2 (Semirings, rings). *A semiring $\mathbb{S}$ is a five-tuple $\langle A, \oplus, \mathbf{0}, \otimes, \mathbf{1} \rangle$ such that $\langle A, \oplus, \mathbf{0} \rangle$ and $\langle A, \otimes, \mathbf{1} \rangle$ are monoids satisfying an annihilation law, i.e. $\forall a \in A.a \otimes \mathbf{0} = \mathbf{0}$, and a distributive law, i.e. $\forall a, b, c \in A.a \otimes (b \oplus c) = (a \otimes b) \oplus (a \otimes c)$. A ring is a six-tuple $\langle A, \oplus, \ominus, \mathbf{0}, \otimes, \mathbf{1} \rangle$ such that $\langle A, \oplus, \mathbf{0}, \otimes, \mathbf{1} \rangle$ is a semiring and $\langle A, \oplus, \ominus, \mathbf{0} \rangle$ a group.*

Remark 1. In the soft constraint tradition [5,15], it suffices to consider semirings, and often $\oplus$ is actually idempotent, hence resulting in a tropical semiring. However, for some of our examples we will resort to rings and, sometimes, to fields, i.e. where there is an additional binary operation $\oslash$ that is defined whenever the second operand is different from $\mathbf{0}$ and such that $a \oslash a = \mathbf{1}$.

Soft Constraints. We now present the notion of soft constraints. It is a straightforward generalization of the one adopted for optimization problems, where $\oplus$ is idempotent. We fix a semiring $\mathbb{S} = \langle S, \oplus, \mathbf{0}, \otimes, \mathbf{1} \rangle$.

Definition 3 (Soft constraints). *Let V be a set of variables and D a finite domain of interpretation. A* (soft) constraint *over $\mathbb{S}$ is a function $c : (V \to D) \to S$ associating a value in S for each variable assignment (or valuation) $\eta : V \to D$ of the variables. C is the set of all the possible constraints that can be built from S, D and V.*

Let $\eta : V \to D$ be a valuation and $c : (V \to D) \to S$ be a constraint. We use $c\eta$ to denote the semiring value obtained when c is applied to η. With $\eta[x := d]$ we denote the valuation η' where $\eta'(y) = \eta(y)$ for all $y \in V \setminus \{x\}$ and $\eta'(x) = d$. Given a set $X \subseteq V$ of variables, we use $\eta \downarrow_X = \eta' \downarrow_X$ to denote the fact that $\eta(x) = \eta'(x)$ for all $x \in X$.

We use $c^{-1} : S \to 2^{V \to D}$ to denote the "inverse" of a constraint c, i.e., $c^{-1}(s)$ is the set of assignments $\varXi = \{\eta \mid c\eta = s\}$. We say that the set $\varXi$ is the set of *solutions* of c with respect to s (in the sense that they map c into a designated value s).

A constraint c often depends only on a subset of the variables in V. Formally, a constraint c depends on the set of variables $X \subseteq V$ if for all valuations η, η' we have $c\eta = c\eta'$ whenever $\eta \downarrow_X = \eta' \downarrow_X$. The smallest such set is called the *support* of c and

denoted as $sv(c)$: It identifies the *relevant* variables of the constraint c. In fact, if $sv(c) = \{x_1, \cdots, x_n\}$, we often write $c[x_1 \mapsto v_1, \cdots, x_n \mapsto v_n] = s$ to denote the fact that $c\eta = s$ for any η that maps each x_i into v_i, since the assignment to other variables is irrelevant.

The projection of a constraint c on a variable x, denoted as $c \Downarrow_x$, is the constraint c' such that $c'\eta = \bigoplus_{d \in D} c(\eta[x := d])$. This means that the variable x is "removed" from the support of c in c' and, in fact, it is easy to check that $sv(c \Downarrow_x) \subseteq sv(c) \setminus \{x\}$.

We call c a *constant* constraint if there exists s such that $c\eta = s$ for all valuation η (and hence, $sv(c) = \emptyset$). With a slight abuse of notation, we shall identify a constant constraint c, where $c\eta = s$, with the semiring value s.

The set of constraints C forms a semiring $\mathbb{C}$, whose structure is lifted from $\mathbb{S}$. More precisely, $(c_1 \star c_2)\eta = c_1\eta \star c_2\eta$ for all $\eta : V \to D$ and $\star \in \{\oplus, \otimes\}$. Combining constraints by the monoidal operators means building a new constraint whose support involves, at most, the variables of the original ones, since it is easily proved that $sv(c_1 \star c_2) \subseteq sv(c_1) \cup sv(c_2)$. The resulting constraint is associated with each tuple of domain values for such variables, which is the element that is obtained by adding/multiplying those associated with the original constraints to the appropriate sub-tuples.

Example 1. Consider the Boolean semiring $\mathbb{B} = \langle \{\mathtt{T}, \mathtt{F}\}, \vee, \mathtt{F}, \wedge, \mathtt{T} \rangle$, a set of variables V and the integer domain $D = [0, 100]$. The resulting set of constraints C whose support contains $\{x, y\} \subseteq V$ includes the constant constraints $\mathtt{T}$ and $\mathtt{F}$, constraints such as $c_1 = \{x \leq 42\}$ (i.e., $c_1[x \mapsto v] = \mathtt{T}$ iff $0 \leq v \leq 42$) and $c_2 = \{y \leq 25\}$, as well as their compositions $c_1 \vee c_2$ and $c_1 \wedge c_2$. As expected, $(c_1 \wedge c_2)[x \mapsto v_x, y \mapsto v_y] = \mathtt{T}$ iff $v_x \leq 42$ and $v_y \leq 25$. Moreover, $(c_1 \wedge c_2) \Downarrow_x = c_1$.

2.2 Opinion Dynamic Models

This section recalls the notion of belief update à la DeGroot [11]. The definitions below are taken from [2,22].

Definition 4 (Influence graph). *An (n-agent) influence graph is a directed graph $G = \langle A, E, I \rangle$ such that A is the set of vertices, $E \subseteq A \times A$ the set of edges, and $I : E \to [0, 1]$ the weight function.*

In the following, the set A of vertices is always given by an interval $\{1, \ldots, n\}$ of integers, and I is extended to $I : A \times A \to [0, 1]$ assuming that $I(i, j) = 0$ if $(i, j) \notin E$.

The vertices in A represent n agents of a community or network. The edges E represent the (direct) influence relationship between these agents, i.e. $(i, j) \in E$ means that agent i influences agent j. The value $I(i, j)$ denotes the strength of the influence, where a higher value means a stronger influence.

As expected, the graph G can be represented as a square matrix M_G of $n = |A|$ rows where $M_{j,i} = I(i, j)$ (i.e., $M_{j,i}$ is the degree of influence of agent i on agent j). Hence, we shall identify the graph G with its corresponding matrix M_G and omit the subindex G in M_G when the influence graph can be deduced from the context. We denote by A_i the set $\{j \mid (j, i) \in E\}$ (i.e., $\{j \mid M_{i,j} \neq 0\}$) of agents with a direct influence over agent i.

At each time unit t, every agent updates their opinion. We use $B^t : A \to [0, 1]$ (that can be seen as a vector of $|A|$ elements) to denote the state of opinion at time t, and B_i^t to denote the opinion of agent i at time-unit t. A DeGroot-like opinion model describes the evolution of agents' opinions about some underlying statement or proposition.

Definition 5 (Opinion model). *An* opinion model *is a triple* $\langle G, B^0, \mu_G \rangle$ *where* $G = \langle A, E, I \rangle$ *is an n-agent influence graph,* $B^0 : A \to [0,1]$ *is the initial state of opinion, and* $\mu_G : [0,1]^n \to [0,1]^n$ *is a state-transition function, called* update function*. For every t, the state of opinion at time* $t+1$ *is given by* $B^{t+1} = \mu_G(B^t)$.

The update function μ_G is dependent on G but can be tuned to model different cognitive biases [2,22], including e.g. *confirmation bias* (where agents are more receptive to opinions that align closely with their own), the *backfire effect* (where agents strengthen their position of disagreement in the presence of opposing views), and *authority bias* (where individuals tend to follow authoritative or influential figures, often to an extreme). In this paper we focus on the following biased update function from [22]

$$B_i^{t+1} = B_i^t + R_i \sum_{j \in A_i} \beta_{i,j}^t M_{i,j}(B_j^t - B_i^t) \tag{1}$$

where $R_i = \frac{1}{\sum_j M_{i,j}}$ if $\sum_j M_{i,j} \neq 0$ and 0 otherwise, and $\beta_{i,j}^t = \beta_{i,j}(B_i^t, B_j^t)$ is a value in $[0,1]$ possibly depending on B_i^t and B_j^t.

A broad class of update functions, generalizing the DeGroot model, can be obtained from Equation (1). Intuitively, updates for an agent i can *weight* disagreements $(B_j - B_i)$ with each one of its neighbors j using functions $\beta_{i,j}$ from the opinion states of i and j to $[0,1]$. These functions are referred to as *(generalized) bias factors*. Notice that the same opinion difference can then be weighted differently by bias factors, depending of current opinions of agents i and j. Thus, intuitively $\beta_{i,j}$ may also be seen as *dynamically changing* the constant influence of j over i, i.e., $M_{i,j}$, depending on their opinions.

Remark 2. If $\beta_{i,j}^t = 1$, Equation (1) corresponds to the update function in the classical DeGroot model. In [2] it is studied the case for $\beta_{i,j}^t = 1 - \mid B_j^t - B_i^t \mid$, which is introduced as the confirmation bias factor of i with respect to j at time t. It is not difficult to prove that in both cases, we always have that $B^{t+1} \in [0,1]^n$.

It should be noted that the formula in Equation (1) can be obtained by manipulating suitable matrices as shown in the following remark.

Remark 3. Consider the vector $\mathbf{1}$ such that $\mathbf{1}_i = 1$, the vector R such that $R_i = \frac{1}{(M\mathbf{1})_i}$ if $(M\mathbf{1})_i \neq 0$ and 0 otherwise, the matrices U^t obtained as the Hadamard product of β^t and M, i.e. $U_{i,j}^t = \beta_{i,j}^t M_{i,j}$, and the diagonal matrices

$$N = diag(R) = \begin{bmatrix} R_1 & 0 & \dots & 0 \\ 0 & R_2 & \dots & 0 \\ \dots & \dots & \dots & \dots \\ 0 & 0 & \dots & R_n \end{bmatrix} \quad V^t = diag(U^t \mathbf{1}) = \begin{bmatrix} \sum_j U_{1,j}^t & 0 & \dots & 0 \\ 0 & \sum_j U_{2,j}^t & \dots & 0 \\ \dots & \dots & \dots & \dots \\ 0 & 0 & \dots & \sum_j U_{n,j}^t \end{bmatrix}$$

It is now easy to see that

$$B_i^{t+1} = B_i^t + N_{i,i} \sum_j U_{i,j}^t (B_j^t - B_i^t) = B_i^t + N_{i,i}((\sum_j U_{i,j}^t B_j^t) - (\sum_j U_{i,j}^t B_i^t)) =$$

$$= B_i^t + N_{i,i}((U^t B^t)_i - (V_{i,i}^t B_i^t))$$

so that, for I the diagonal identity matrix, we have

$$B^{t+1} = B^t + N(U^t B^t - V^t B^t) = (I + N(U^t - V^t))B^t \tag{2}$$

where $(U^t - V^t)\mathbf{1} = 0$, and hence, $(I + N(U^t - V^t))\mathbf{1} = \mathbf{1}$. In other terms, $I + N(U^t - V^t)$ is a *row stochastic* matrix, where the sum of the elements in each row is 1.

In the case of $\beta_{i,j} = 1$, we have that $U^t = M$ and $V^t = diag(M\mathbf{1})$, so that $NV = I$ and $B^{t+1} = NMB^t$. If M is also assumed to be a row stochastic matrix, i.e. it satisfies $M\mathbf{1} = \mathbf{1}$, then $N = I$ and thus $B^{t+1} = MB^t$, as expected in the standard deGroot model. $\qquad\square$

Given a matrix M representing the influence graph G of an opinion model, we shall use M^* to denote the limit of M^t for $t \to \infty$. If M is row stochastic and, furthermore, *strongly connected* (i.e. the influence graph it represents is strongly connected) and *aperiodic* (i.e. the greatest common divisor of the lengths of its cycles is one) [16], the limit M^* exists and the rows of M^* are all the same. In this case, for all pair of agents i and j and initial beliefs B, $(M^*B)_i = (M^*B)_j$. Thus, assuming $\beta_{i,j} = 1$, so that $B^{t+1} = MB^t$ as in the DeGroot model, we obtain that, in the limit, the opinion of the agents converges to the same value, thus reaching a *consensus* independently from their initial beliefs. A simple and sufficient condition for M to be aperiodic is by checking that there exists an agent i such that $M_{i,i} > 0$. An agent where $M_{i,i} = 0$ can be seen as a *puppet* that just incorporates the opinions of others.

Theorem 1 (Consensus in the DeGroot model [11]). *Let M be a row stochastic matrix of values in $[0,1]$. If M is strongly connected and aperiodic, then $\lim_{t\to\infty} M^t$ exists and all its rows are equal.*

A consensus result for more general bias factors $\beta_{i,j}$ under some suitable conditions (including continuity) can be found in [22].

3   Constraint Opinion Models

Opinion models, as shown in the previous section, represent opinions and influences as a real number in the interval $[0,1]$. This section generalizes this idea and proposes a *constraint opinion model*, where the opinions and the influences of the agents are represented as (soft) constraints in a given semiring.

Recall from Section 2.1 that given a set of variables V, a finite domain D, and a semiring $\mathbb{S}$, the set of constraints built from $\mathbb{S}$, D and V is denoted C. In fact, we have a semiring $\mathbb{C}$ of constraints, with carrier C, whose structure lifted from $\mathbb{S}$.

Definition 6 (Constraint influence graph). *Let $\mathbb{S}$ be a semiring and C the constraints built from $\mathbb{S}$, a set of variables V, and a finite domain D. Let $A = \{1, \cdots, n\}$ be a set of agents. An (n-agents) constraint influence graph is a square matrix M of dimension n with values in C, i.e. $M : A \times A \to C$.*

An element $c_{i,j}$ of M (row i, column j) represents how agent j influences agent i. As shown below, this influence can be a constant or an arbitrary constraint, reflecting scenarios where agents impose specific restrictions on the way they are influenced.

Definition 7 (Constraint Opinion Model). *A* constraint opinion model *is a triple* $\langle M, B^0, \mu^t \rangle$ *where M is an n-agent constraint influence graph, $B^0 : A \to C$ is the initial state of opinion, and $\mu^t : C^n \to C^n$ is the update function at time t, so that the state of opinion at time $t+1$ is given by $B^{t+1} = \mu^t(B^t)$.*

As shown in the forthcoming sections, representing the opinion of an agent at time t as a constraint B_i^t, widens the spectrum of situations that can be modeled in systems of agents updating their beliefs. In particular, an opinion model (Definition 5) is an instance of a constraint opinion model where the set of variables V is $\{p\}$, the domain of the variable p is $D = \{0,1\}$, the semiring is given by the positive real numbers $\langle \mathbb{R}^+, +, 0, \times, 1 \rangle$. Any opinion c is required to satisfy $\sum_\eta c\eta = 1$ (e.g., if $c[p \mapsto 0] = 0.3$, then $c[p \mapsto 1] = 0.7$), all the elements in the matrix (influence graph) M are constants, and the update function has type $\mu^t : [0,1]^n \to [0,1]^n$.

In the following examples, it is assumed that μ^t is given by the following matrix multiplication equation:

$$B^{t+1} = \mu^t(B^t) = MB^t = M^t B^0. \tag{3}$$

Example 2 (Opinion Models in $\mathbb{B}$). Consider again the semiring $\mathbb{B}$ and the constraints $c_1 = \{x \le 42\}$ and and $c_2 = \{y \le 25\}$ in Example 1. Let $d_1 = c_1 \wedge c_2 = \{x \le 42, y \le 25\}$, $d_2 = c_3 \vee c_4$ where $c_3 = \{x \ge 15\}$ and $c_4 = \{y \ge 66\}$, and let $B^0 = [d_1 \ d_2]$ be the initial set of opinions. Consider also the following constraint influence graphs

$$M_1 = \begin{bmatrix} \text{T} & \text{F} \\ \text{T} & \text{T} \end{bmatrix} \qquad M_2 = \begin{bmatrix} \text{T} & \text{F} \\ \text{F} & \text{T} \end{bmatrix} \qquad M_3 = \begin{bmatrix} \text{F} & \text{T} \\ \text{T} & \text{F} \end{bmatrix} \qquad M_4 = \begin{bmatrix} \text{T} & y \le 20 \\ x \ge 10 & \text{T} \end{bmatrix}$$

The matrix M_1 is idempotent (i.e. $M_1 M_1 = M_1$) and represents the situation where Agent 1 is not influenced at all by Agent 2 while, instead, Agent 2 accepts all the information from Agent 1. In this scenario, we have $M_1 B^0 = B^1 = [d_1 \ d_1 \vee d_2]$, and $M_1 B^1 = B^1$. This means that, after one interaction, the system stabilizes in an opinion where Agent 1 does not change its initial opinion d_1, and Agent 2 considers also possible the values for x and y according to the opinion d_1 of Agent 1.

M_2 is the identity matrix I, and it represents a situation where neither agent is influenced by the other. Hence, $M_2 B = B$ for any choice of B. Instead, M_3 is involutory (i.e. $M_3 M_3 = I$) and thus the system never stabilizes, alternating between $[d_1 \ d_2]$ and $[d_2 \ d_1]$.

The constraint influence graph M_4 is more interesting, since some of the influences are not constants. Consider the element $(M_4)_{2,1} = \{x \ge 10\}$: Agent 2 is influenced by Agent 1 only to the point that it accepts that x might also be bigger than 10 (but not smaller than that). We thus have $M_4 B^0 = [d_1' \ d_2']$ with

$$d_1' = (c_1 \wedge c_2) \vee (\{y \le 20\} \wedge (c_3 \vee c_4)) = (c_1 \wedge c_2) \vee \{x \ge 15, y \le 20\}$$
$$d_2' = (\{x \ge 10\} \wedge c_1 \wedge c_2) \vee (c_3 \vee c_4) = \{10 \le x \le 42, y \le 25\} \vee (c_3 \vee c_4)$$

The use of constraints in M_4 allows us to represent a *core belief* [19,18] for Agent 2, as it "filters" (part of) the opinion of Agent 1 when it is not consistent with the limits she imposes to update her opinions. In this case, after interaction, Agent 2 still cannot believe that $x = 9$, a scenario that Agent 1 considers plausible. This is possible due to

the flexibility of using constraints on the influence graph (and not only on opinions): Agents can impose different conditions on how they are influenced by other agents.

Example 3 (Opinion Models in $\mathbb{R}^+$). Consider the 2-agent influence graph below

$$M = \begin{bmatrix} 1 & 0 \\ 0.8 & 0.2 \end{bmatrix}$$

Agents 1 and 2 discuss about a proposition p: Agent 1 tends to believe that p is not the case while Agent 2 is more positive about p. We represent this situation with the constraints $c_1 = \{0 \mapsto 0.3, 1 \mapsto 0.7\}$ and $c_2 = \{0 \mapsto 0.6, 1 \mapsto 0.4\}$, respectively. According to M, Agent 1 accepts no influence, while Agent 2 does. Assume that the update function μ^t is matrix multiplication. We have $MB^0 = [c'_1 \; c'_2]$ with $c'_1 = c_1$ and $c'_2 = \{0 \mapsto 0.4, 1 \mapsto 0.6\}$. Since $M_{1,2} = 0$, Agent 1 does not change her opinion. Due to the influence of Agent 1, the opinion of Agent 2 progressively converges to c_1.

3.1   On the update function

Consider again the calculations in Remark 3. Note that they can be mimicked in any field, so that given the Hadamard product $U^t = \beta^t \otimes M$ we have

$$B^{t+1} = (I \oplus N(U^t \ominus V^t))B^t \tag{4}$$

Note that in this general case we also have $(U^t \ominus V^t)\mathbf{1} = 0$, so that $(I \oplus N(U^t \ominus V^t))\mathbf{1} = \mathbf{1}$, and hence $I \oplus N(U^t \ominus V^t)$ is a row stochastic matrix. If, moreover, M is also row stochastic, we have $B^{t+1} = (I \oplus U^t \ominus V^t)B^t$, which is a valid equation for any ring, since the definition of the matrices does not involve the division operator, if β^t does not.

Let us now consider the two examples above: They adopt the standard DeGroot model, i.e. such that $B^{t+1} = MB^t$, which is a valid equation for any semiring. All the elements of the first three matrices in Example 2 are constant constraints. Also, the four matrices are strongly connected and row stochastic, since in each row the sum of the constraints gives the constant constraint always returning T. In particular, for M_4 we have that $\mathsf{T} \vee y \leq 20 = x \geq 10 \vee \mathsf{T} = \mathsf{T}$. However, M_3 is not aperiodic, since the only cycle has length 2. The matrix in Example 3 contains only constant constraints and it is row stochastic, strongly connected, and aperiodic, since each agent has a cycle.

We now start investigating the model on semirings. The following definition makes precise the idea of a set of agents reaching a *consensus* (as the two agents in Example 3) and of a set of agents agreeing in the limit about a particular valuation for the variables (see Example 7 in the following section).

Definition 8 (Consensus). *Let $\langle M, B^0, \mu^t \rangle$ be a constraint opinion model where $M : A \times A \to C$ is a constraint influence graph. We say that the set of agents A converges to an opinion $c \in C$ whenever for each $i \in A$, $\lim_{t \to \infty} B_i^t = c$. We say that the set of agents A converges to a consensus whenever each agent in A converges to the same opinion. Given a valuation η, we say that the agents converge to a consensus about η whenever there exists $s \in S$ such that for each $i \in A$, $\lim_{t \to \infty} B_i^t \eta = s$.*

We now move a step further and we consider the set C_p of what we call *probability* constraints, that is, $c \in C_p$ if $\oplus_\eta c \eta = 1$.

Proposition 1. *Let M be a row stochastic matrix of constant constraints and B a vector whose elements belong to C_p. Then also the elements of the vector MB belong to C_p.*

4   Sharing and Discussing about Partial Information

This section illustrates how constraint opinion models provide an extra expressiveness for representing scenarios involving agents sharing opinions. We explore situations that cannot be modeled using the traditional DeGroot framework, such as agents discussing preferences and exchanging partial information (Section 4.1), modeling conditional preferences (Section 4.2), and capturing agents' beliefs about one another (Section 4.3). In all the examples reported here, the update function is given as in Eq.3.

4.1   Preferences and Partial Information

We consider agents engaging in discussions about a given proposition p. As noted in Section 3, when the constraint opinion model is instantiated to represent the classic De-Groot model, the domain of p is restricted to $\{0,1\}$. In the section, instead, we consider a different domain for p allowing us to represent situations where agents give different *preferences* to p, or when they only have *partial information* about p. This will be useful to model opinions that cannot be represented in the DeGroot model (as a single real number in the interval $[0,1]$).

Consider a variable p with finite domain $D = \{1, \cdots, 5\}$, where 1 means "very bad" and 5 means "very good" as in a Likert scale, or alternatively, "extreme-left" and "extreme-right", "strongly disagree" and "strongly agree", etc. Let $a, b \in D$ such that $a \leq b$ and define the constraint $\iota_{(a,b)}$ as $\iota_{(a,b)}(v) = \frac{1}{1+(b-a)}$ if $v \in [a,b]$ and $\iota_{(a,b)}(v) = 0$ otherwise. The constraint $\iota_{(a,b)}$ represents an opinion where the agent assigns equal preferences (different from 0) when the value of p is in the interval $[a,b]$ and 0 otherwise. Note that for all a, b such that $1 \leq a \leq b \leq 5$, $\iota_{(a,b)} \in C_p$ (i.e. $\oplus_\eta \iota_{(a,b)} \eta = 1$).

Consider the following inference graphs and initial opinion

$$M_1 = \begin{bmatrix} 0.9 & 0.1 \\ 0.8 & 0.2 \end{bmatrix} \quad M_2 = \begin{bmatrix} 0.1 & 0.9 \\ 0.2 & 0.8 \end{bmatrix} \quad M_3 = \begin{bmatrix} 0.5 & 0.5 \\ 0.5 & 0.5 \end{bmatrix} \quad B_0 = \begin{bmatrix} \iota_{(1,2)} & \iota_{(1,5)} \end{bmatrix}$$

Regarding the initial opinions, Agent 1 is somewhat negative about p but *uncertain* whether to assess p as "very bad" or merely "bad." In contrast, Agent 2 has no specific opinion about p and thus assigns an equal preference 0.2 to all possible values of p. It is worth noticing that the partial information conveyed by these opinions cannot be accurately expressed in the DeGroot model as a single value within the interval $[0,1]$.

When the influence graph M_1 is considered, the system converges to a consensus where the two agents assign a preference of approximately 0.467 when $p \in \{1,2\}$ and 0.022 otherwise. In this scenario, Agent 1 has a stronger influence on Agent 2, hence the agents tend to agree that p is "very bad" or "bad". In the case of M_2, the system

(a) Initial opinions given in Example 4   (b) Initial opinions given in Example 6
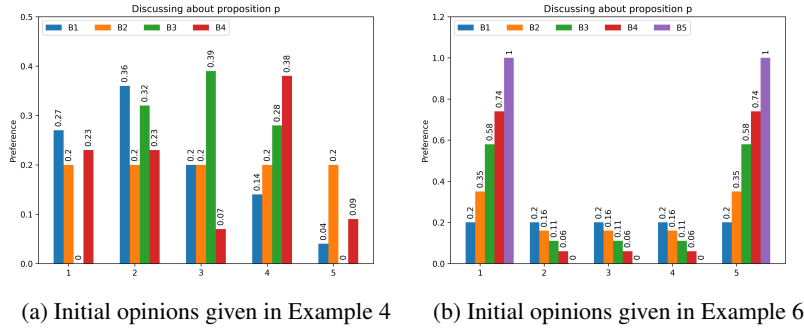
Fig. 1: Preferences in the consensus when agents discuss a proposition p.

converges to a preference of approximately 0.255 when $p \in \{1, 2\}$ and 0.163 otherwise: since Agent 2 has a strong influence on Agent 1, in the end Agent 2 considers more plausible that p is not "too bad". Finally, if we consider M_3, where the agents have the same level of influence, the system converges to a preference of approximately 0.35 when $p \in \{1, 2\}$ and 0.1 otherwise.

Example 4. Let $\iota_{(a,b)}$ be as above and consider the following influence graph and vectors of initial opinions, where agents are not completely sure how to assess p (*partial information*)

$$
M = \begin{bmatrix} 0.2 \ 0.3 \ 0.4 \ 0.1 \\ 0.3 \ 0.1 \ 0.2 \ 0.4 \\ 0.5 \ 0.1 \ 0.2 \ 0.2 \\ 0.3 \ 0.3 \ 0.3 \ 0.1 \end{bmatrix} \qquad \begin{aligned} B_1 &= \begin{bmatrix} \iota_{(1,2)} \ \iota_{(1,5)} \ \iota_{(2,4)} \ \iota_{(1,3)} \end{bmatrix} \\ B_2 &= \begin{bmatrix} \iota_{(1,5)} \ \iota_{(1,5)} \ \iota_{(1,5)} \ \iota_{(1,5)} \end{bmatrix} \\ B_3 &= \begin{bmatrix} \iota_{(2,2)} \ \iota_{(3,3)} \ \iota_{(4,4)} \ \iota_{(3,3)} \end{bmatrix} \\ B_4 &= \begin{bmatrix} \iota_{(1,2)} \ \iota_{(1,3)} \ \iota_{(4,4)} \ \iota_{(4,5)} \end{bmatrix} \end{aligned}
$$

The consensus values obtained by starting with each of the initial opinions B_i above are shown in Figure 1a. For B_1, the agents generally assign a negative score to p, with a stronger preference for $p = 2$ (a value present in all the initial beliefs). In B_2, the agents equally prefer all possible values of p, achieving a consensus where no specific value is favored over others. For B_3, the agents are more certain about the value of p, with two of them believing that $p = 3$. Consequently, $p = 3$ becomes the most preferred value in the consensus, followed by $p = 2$. In this scenario, note also that none of agents believe that $p = 1$ or $p = 5$ and those values receive a preference of 0 in the consensus. In B_4, two agents are inclined to assign a negative score to p, while the other two are more positive. Given this particular influence graph, the system converges to a situation where $p = 4$ receives the highest preference. □

The examples above consider opinions to be elements in C_p. Now we consider scenarios where agent's opinions do not adhere to that restriction.

Preferance for Extreme Positions. The following example considers agents that prefer any extreme option than a moderate one. This kind of preference appears for example in risky decision-making, where an agent might prefer high-risk, high-reward options over safe middle-ground ones.

Example 5. Consider the following constraint

$$c(x,y) \stackrel{\text{def}}{=} \lambda p. \begin{cases} x & \text{if } p = 1 \text{ or } p = 5 \\ y & \text{otherwise} \end{cases}$$

Intuitively, the opinion $c_e = c(1.0, 0.0)$ represents the situation where the agent "likes the extremes" (regardless of which one), while the constraint $c_c = c(0.0, 1.0)$ represents the opinion of an agent that prefers more moderate positions.

Consider the following situation (note that M satisfies the conditions in Theorem 1)

$$M = \begin{bmatrix} 0.3 & 0.7 \\ 0.4 & 0.6 \end{bmatrix} \quad \begin{matrix} B_1 = \begin{bmatrix} c_e & d \end{bmatrix} \\ B_2 = \begin{bmatrix} c_e & c_c \end{bmatrix} \end{matrix}$$

where $d[p \mapsto 2] = 0.2$, $d[p \mapsto 3] = 0.3$ and $d[p \mapsto v] = 0.0$ otherwise. Starting from the initial beliefs B_1 and B_2 above, the system converges to the following opinions c_1 and c_2 respectively

$$c_1 = \{1 \mapsto 0.36, 2 \mapsto 0.13, 3 \mapsto 0.19, 4 \mapsto 0.00, 5 \mapsto 0.36\}$$
$$c_2 = \{1 \mapsto 0.36, 2 \mapsto 0.64, 3 \mapsto 0.64, 4 \mapsto 0.64, 5 \mapsto 0.36\}$$

In c_1, the agents assign higher preferences (but less than the original 1.0 for Agent 1) to the extreme positions. Since Agent 2 has a stronger influence on Agent 1, the opinion c_2 assigns higher preferences to the moderate positions than the extreme ones.   □

Example 6. Consider the influence graph M in Example 4 and the initial opinions

$$B_1 = \begin{bmatrix} \iota_{(1,5)} & \iota_{(1,5)} & \iota_{(1,5)} & \iota_{(1,5)} \end{bmatrix} \qquad B_2 = \begin{bmatrix} \iota_{(1,5)} & \iota_{(1,5)} & \iota_{(1,5)} & c_e \end{bmatrix}$$
$$B_3 = \begin{bmatrix} \iota_{(1,5)} & \iota_{(1,5)} & c_e, c_e \end{bmatrix} \qquad B_4 = \begin{bmatrix} \iota_{(1,5)} & c_e & c_e & c_e \end{bmatrix}$$
$$B_5 = \begin{bmatrix} c_e & c_e & c_e & c_e \end{bmatrix}$$

Starting with B_1, where none of the agents have a formed opinion about p (constraint $\iota_{(1,5)}$), we progressively add agents with "extreme" point of views (constraint c_e). For all these configurations, Figure 1b shows the constraints obtained in the consensus. As expected, B_1 (respectively, B_5) is already a consensus where no value of p is more preferred than other (resp., the extreme values are equaled preferred and the moderate ones are not considered). For this particular configuration of M, with 2 "extreme" agents (opinion B_3), there is a clear tendency to prefer the valuations $p = 1$ and $p = 5$.   □

Discussing several topics. A constraint opinion model with one variable taking values from a finite domain D can be also interpreted as agents discussing about different topics simultaneously as the following example shows.

Example 7. Let $D = \{1, 2, 3\}$ and consider the following definition

$$c_m(x_1, x_2, x_3) \stackrel{\text{def}}{=} \lambda p. \begin{cases} x_1 & \text{if } p = 1 \\ x_2 & \text{if } p = 2 \\ x_3 & \text{otherwise} \end{cases}$$

The constraint $d = c_m(v_1, v_2, v_3)$ represents the opinion of an agent about three different propositions, where the opinion about proposition p_i is $d[p \mapsto i]$. As an example, consider the situation

$$M = \begin{bmatrix} 0.3 & 0.7 \\ 0.5 & 0.5 \end{bmatrix} \qquad B_0 = [\, c_m(0.3, 0.7, 0.1) \quad c_m(0.4, 0.4, 0.4) \,]$$

Agent 1 is positive about the second proposition and tends to be more negative about the other two propositions. In this case, the system converges to the following situation:

$$c = \{1 \mapsto 0.36, 2 \mapsto 0.52, 3 \mapsto 0.26\}$$

When non-constant constraints are considered in the influence graph, it is possible to represent the situation where the influence of the agents depends on the proposition being discussed. For example, consider the initial opinion B_0 above and the following influence graph

$$M = \begin{bmatrix} c_m(0.3, 0.2, 0.1) & c_m(0.7, 0.8, 0.9) \\ c_m(0.5, 0.1, 0.8) & c_m(0.5, 0.9, 0.2) \end{bmatrix}$$

Note, for instance, that Agent 1 exerts less influence over Agent 2 when discussing about the second proposition (0.1) than when discussing the third proposition (0.8). For this configuration, the agents converge to the following opinion

$$c = \{1 \mapsto 0.36, 2 \mapsto 0.43, 3 \mapsto 0.26\}$$

$\square$

In the example above, the influence graph is not a matrix of constant constraints; therefore, the consensus theorem on reals does not directly apply. However, in this specific setting, where the influences of different topics are independent, we can naturally extend Theorem 1 as follows.

Proposition 2. *Consider a constraint opinion model on $\mathbb{R}^+$ with one variable p with finite domain D and an influence graph M. If for all $d \in D$ the matrix of constant constraints $M\{p \mapsto d\}$ (the result of applying η to each element $c_{i,j}$ in M) is row stochastic, strongly connected and aperiodic, then $\lim_{t \to \infty} M^t$ exists and all its rows are equal.*

4.2   Partial Information and Conditional Opinions

Example 2 showed that using crisp constraints (semiring $\mathbb{B}$), it is possible to represent the situation where agents share partial information about the actual value of a given variable. Moreover, using constraints as influences, it is possible to "filter" the information coming from other agents. In this section we show how this idea can be further generalised when $\mathbb{B}$ is replaced with $\mathbb{R}^+$.

Consider the situation where four agents are trying to decide the number of members a selection committee must have. We represent that decision with a variable x. Agents

have different opinions about the actual value of x, and such opinions can be naturally expressed as constraints. For instance, the initial belief could be given by

$$o_1 : 4 \leq x \leq 6 \qquad o_2 : 5 \leq x \leq 10 \qquad o_3 : 6 \leq x \leq 7 \qquad o_4 : 1 \leq x \leq 4$$

The opinions above express different degrees of uncertainty, which are reflected on the number of possible values o_i allows. The initial belief of each agent is a constraint c_i such that $c_i[x \mapsto v] = 1$ if $o_i[v/x]$ is true and $c_i[x \mapsto v] = 0$ otherwise. For instance, the constraint $\{4 \leq x \leq 6\}$ is the function that maps to 1 the valuations that map x to a value $v \in \{4, 5, 6\}$.

The four agents need to take a decision and they define how they will be influenced by the others' opinions, for instance

$$M = \begin{bmatrix} 0.3 & 0.2 & 0.3 & 0.2 \\ 0.2 & 0.3 & 0.1 & 0.4 \\ 0.2 & 0.4 & 0.2 & 0.2 \\ 0.1 & 0.1 & 0.5 & 0.3 \end{bmatrix}$$

After some iterations, the system converges to the following constraint:

$$\begin{array}{ll} c[x \mapsto v] = 0.28 \text{ if } 1 \leq v \leq 3 & c[x \mapsto 4] = 0.46 \\ c[x \mapsto 5] = 0.44 & c[x \mapsto 6] = 0.72 \\ c[x \mapsto 7] = 0.53 & c[x \mapsto v] = 0.25 \text{ if } 8 \leq v \leq 10 \\ c[x \mapsto v] = 0 \text{ otherwise} & \end{array}$$

Hence, a selection committee of 6 people is the best solution these agents can find that "better" satifies their initial beliefs and influences.

Now suppose that the agents must also decide whether the selection committee must include external members or not. Such a decision is modeled with a second variable y with domain $\{0, 1\}$. Using constraints, it is natural to represent opinions such as "with external members, I would prefer a committee of size $5 \leq x \leq 6$ but, without them, I would prefer a committee of size $3 \leq x \leq 5$". This statement can be modeled as the implication $y = 1 \Rightarrow 5 \leq x \leq 6$ (or $y = 0 \vee 5 \leq x \leq 6$) and $y = 0 \Rightarrow 3 \leq x \leq 5$. We shall use $c_s(5, 6, 3, 5)$ to represent such an opinion where c_s is defined as

$$c_s(a_0, b_0, a_1, b_1) \stackrel{\text{def}}{=} \lambda x \, y.((\{y = 1\} + \{a_0 \leq x \leq b_0\}) \times (\{y = 0\} + \{a_1 \leq x \leq b_1\}))$$

Consider also the definition

$$c'_s(s, a, b) \stackrel{\text{def}}{=} \lambda x \, y.(\{y = s\} \times \{a \leq x \leq b\})$$

The constraint $c'_s(1, 4, 6)$ represents the opinion "the committee must be formed by a number of $4 \leq x \leq 6$ people and it must include external members".

Figure 2a shows the consensus c reached with the influence graph M above, when the initial vector of beliefs is

$$B_0 = [\, c_s(5, 6, 3, 5) \quad c_s(8, 10, 7, 9) \quad c_s(5, 7, 2, 4) \quad c'_s(1, 4, 6) \,]$$

The green bars represent the opinion $c \Downarrow_x$ (the projection of c onto the variable x). This projection captures the agents' preferences when disregarding their opinions about

(a) Preferences in the Consensus
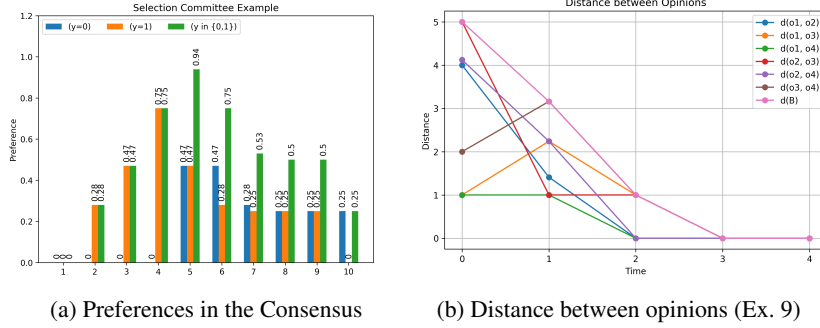
(b) Distance between opinions (Ex. 9)

Fig. 2: Preferences in the consensus when agents discuss the size of a committee.

the inclusion of external members. Similarly, the constraint $d = c \Downarrow_y$ represents the preferences of the agents regarding only the value of y. In this example, $d[y \mapsto 0] = 1.97$ and $d[y \mapsto 1] = 3.0$. Thus, if we select the best assignment (constraint c, blue and orange bars), the optimal choice is to form a selection committee of 4 members, including external ones. However, if we ignore the decision about external members ($c \Downarrow_x$, green bars), the preferred option is a selection committee of 5 members. Finally, the constraint $c \Downarrow_y$ indicates that the agents favor the inclusion of external members.

4.3   Beliefs

Hitherto we have assumed that the opinion of an agent $a \in A$ is a constraint of type $c_a : (X \to D) \to S$. In this section we aim to be more nuanced, and to be able to model the beliefs of agents about themselves and about the other agents. To this end, a constraint should be of the form $c_a : A \to ((X \to D) \to S)$. Consider e.g. the semiring $\mathbb{B}$ and the agents $A = \{a, b\}$. The "constraint" c_a such that $c_a(a) = \{x > 42\}$ and $c_a(b) = \{x \leq 20\}$ represents the situation where a thinks that $x > 42$ and that b believes that $x \leq 20$.

Now, note that $A \to ((X \to D) \to S)$ is equivalent to $(A \times (X \to D)) \to S$ and $A \times (X \to D)$ is equivalent to $(\{\bullet\} \to A) \times (X \to D)$, which in turn is equivalent to the *typed* functions $(\{\bullet\} \uplus X) \to (A \uplus D)$. Finally, note that the results of the previous sections could be rephrased for constraints over typed valuations, even if we restrained from doing so for the sake of simplicity. For the rest of this section we then consider valuations for typed variables $\{\bullet\} \uplus X$ and disjoint domains $A \uplus D$. Also, we represent a constraint c over such valuations as $\bigoplus_{a \in A} c^a$ for $c^a \eta = c\eta$ if $\eta(\bullet) = a$ and 0 otherwise.

For the influence graph, the elements $M_{i,j}$ of the matrix are thus constraints that can be also represented as $M_{i,j} = \bigoplus_a M_{i,j}^a$, so that $M = \bigoplus_a M^a$. Since also a vector B can be decomposed as $\bigoplus_a B^a$, we have that $MB = \bigoplus_a M^a B^a$ We can further impose the restriction that each component matrix is row stochastic, i.e. $M^a \mathbf{1} = \mathbf{1}$. In this way, the results and definitions in Section 3.1 can be recast to the case of beliefs.

Similar to Example 2, consider the semiring $\mathbb{B}$, a set of two agents $A = \{a, b\}$ and the following constraints and influence graphs

$$c_a = \{x > 42, y \leq 30\}^a \oplus \{x \leq 20\}^b \qquad c_b = \{x > 10, y = 20\}^b$$

$$M = \begin{bmatrix} T^a \oplus F^b & F^a \oplus T^b \\ T^a \oplus T^b & T^a \oplus T^b \end{bmatrix} \qquad\qquad N = \begin{bmatrix} T & F \\ T & T \end{bmatrix} = \begin{bmatrix} T^a \oplus T^b & F^a \oplus F^b \\ T^a \oplus T^b & T^a \oplus T^b \end{bmatrix}$$

Agent a beliefs that $x > 42$ and $y \leq 30$. Moreover, Agent a thinks that Agent b's opinion is $x \leq 20$. On the other side, Agent b thinks that $x > 10$ and $y = 20$, and Agent b does not have any opinion about the values Agent a considers possible for x and y.

In M, note that $M_{1,1}^b = F$ and $M_{1,2}^b = T$. This means that a is willing to scrap its belief about b and to accept the opinion of Agent b about itself. If we compose $M[\begin{array}{cc} c_a & c_b \end{array}]$ we obtain $[\begin{array}{cc} d_a & d_b \end{array}]$ with

$$d_a = c_a^a \oplus c_b^b = \{x > 42, y \leq 30\}^a \oplus \{x > 10, y = 20\}^b$$

$$d_b = c_a \oplus c_b = \{x > 42, y \leq 30\}^a \oplus \{x \leq 20\}^b \oplus \{x > 20, y = 20\}^b$$

where the latter equality is due to the fact that

$$\{x \leq 20\} \oplus \{x > 10, y = 20\} = \{x \leq 20\} \oplus \{x > 20, y = 20\}$$

Finally, composing $N[\begin{array}{cc} c_a & c_b \end{array}]$ we obtain $[\begin{array}{cc} d_a & d_b \end{array}]$ with $d_a = c_a$ and $d_b = c_a \oplus c_b$.

5   Polarization Measures

Polarization measures aim to quantify how divided a set of agents is concerning their opinions [13]. A key element in this assessment is the ground distance between opinions, which serves as a way to determine whether disagreements are minor or extreme. When opinions are exact values represented as real numbers, as assumed in previous models for opinion dynamics, the Euclidean distance and the absolute difference are natural choices. Nevertheless, in the model proposed here, opinions are not precisely known but are instead represented by constraints. Hence, defining a ground distance is more complex, as we no longer compare single points but uncertainty regions.

In this section we introduce a polarization measure for constraint opinion models, aimed at quantifying the "divergence" between two constraints. The key idea is to measure this divergence based on the distance between their respective solution sets. Remember that, given a semiring value s, $c^{-1}(s)$ is the set of assignments an agent considers possible, or at least, those that are consistent with a "preference level" s. By evaluating this distance, we gain insight into the degree of alignment or divergence between agents' opinions.

We fix a set of variables X, a domain of interpretation D, and a semiring $\mathbb{S} = \langle S, \oplus, \mathbf{0}, \otimes, \mathbf{1} \rangle$. At first sight, it seems natural to define the distance between constraints by assuming that $\mathbb{S}$ is a metric space, equipped with a distance $\delta : S \times S \to \mathbb{R}^+$. However, this allows to capture *only* the distance between two particular assignments but it does not tell us much about *all* the possible assignments the agents consider plausible.

Moreover, in the case of the Boolean semiring, this measure will be the coarsest possible: the distance is 0 if $c_1\eta = c_2\eta$, or a value $\delta(\mathtt{T},\mathtt{F}) = \delta(\mathtt{F},\mathtt{T})$ otherwise. The distance proposed here, instead, assumes that D is a metric space that we lift to a metric space on valuations under the assumption of a finite support for variables.

If D is a metric space, equipped with a distance $\delta : D \times D \to \mathbb{R}^+$, we can define a new metric space D^n for any $n \in \mathbb{N}$, assuming the existence of a norm function $\|\cdot\| : \mathbb{R}^n \to \mathbb{R}^+$, so that we have

$$\delta((x_1,\ldots,x_n),(y_1,\ldots,y_n)) = \|\delta(x_1,y_1),\ldots,\delta(x_n,y_n)\|$$

The most used norm in $\mathbb{R}^n$ is the Euclidean norm L^2, so that we have

$$\|(z_1,\ldots,z_n)\| = \sqrt{z_1^2 + \ldots + z_n^2}$$

Consider for instance the examples in Section 4, where D is a finite subset of $\mathbb{N}$. We can consider the usual distance for D, namely $\delta(x,y) = |x-y|$. Adopting the L^2 norm for $\mathbb{R}^n$ results in the Euclidean distance on $\mathbb{N}^n$ for every n.

Remark 4. In the rest of this section, we restrict constraints to be built over a finite set V of variables, and we assume that the domain D is a metric space, lifted to $D^{|V|}$.

Given a semiring value s, the distance between two constraints c and d will be the distance between the sets of assignments $c^{-1}(s)$ and $d^{-1}(s)$, thus comparing the "solutions" of the opinions c and d with respect to a preference level s. Since $V \to D$ is a metric space, we can compare two sets of assignments as follows.

Definition 9 (Hausdorff Distance). *Let $\mathbb{M}$ be a metric space. The distance between an element of the metric space and a subset of elements of $\mathbb{M}$ is defined as $\delta(v,B) = \inf_{w\in B}\delta(v,w)$. The forward distance between two subsets of $\mathbb{M}$ is defined as $\overrightarrow{\delta}(A,B) = \sup_{v\in A}\delta(v,B)$. For two non-empty sets A and B, the Hausdorff distance is defined as*

$$\delta_H(A,B) = max\left(\overrightarrow{\delta}(A,B), \overrightarrow{\delta}(B,A)\right)$$

Moreover, $\delta_H(A,\emptyset) = \delta_H(\emptyset,A) = \infty$ and $\delta_H(\emptyset,\emptyset) = 0$.

Intuitively, the distance $\delta_H(A,B)$ is the greatest of all the distances from a point in the set A to the closest point in the set B. Notice that this guarantees that any element of one set is within a distance of *at most* $\delta_H(A,B)$ of some element of the other set.

Definition 10 (Distance between opinions). *Let $s \in S$ and c,d be two constraints. We define $\delta_s(c,d) = \delta_H(c^{-1}(s),d^{-1}(s))$ and*

$$\delta(c,d) = max\{\delta_s(c,d) \mid s \in S\}$$

Both definitions are well-given since V is finite, hence the images of c and d are so.

These definitions can be generalized to a finite set O of constraints, so that e.g. $\delta_s(O) = max\{\delta_s(c_i,c_j) \mid c_i,c_j \in O\}$. Such a measure can be interpreted as an indicator of *polarization* in the opinions of a set of agents. A higher value of $\delta_s(O)$ corresponds to a greater distance between the two most *antagonistic* agents. Moreover, small values of $\delta_s(O)$ indicate that agents tend to consider as plausible the same set of solutions.

If the semiring $\mathbb{S}$ is equipped with an order $\leq$, we can define $c_{\geq}^{-1}(s)$ as the set of valuations $\{\eta \mid s \leq c\eta\}$, i.e., the valuations that assigns a value s' at least as "good" as s. Accordingly, we define $\delta_{\geq s}(c,d) = \delta_H(c_{\geq}^{-1}(s), d_{\geq}^{-1}(s))$ and, for a set of opinions O, $\delta_{\geq s}(O) = max\{\delta_{\geq s}(c_i,c_j) \mid c_i,c_j \in O\}$.

Example 8. Consider the constraints in Example 2, choosing $V = \{x,y\}$ and let

$$A = (c_1 \wedge c_2)^{-1}(\mathrm{T}) = \{x \leq 42\} \cap \{y \leq 25\}$$

$$B = (c_3 \vee c_4)^{-1}(\mathrm{T}) = \{x \geq 15\} \cup \{y \geq 66\}$$

where A (respectively B) is the set of valuations that make the constraint $c_1 \wedge c_2$ (respectively $c_3 \vee c_4$) true. In this example, $D = \{0,\dots,100\}$ and we use D^2 as metric space, adopting the Euclidean distance.

The intersection $A \cap B$ is $\{15 \leq x \leq 42\} \cap \{y \leq 25\}$. Concerning the remaining elements of A, i.e. those satisfying $x < 15$, the minimal distance with respect to B is given by $15 - x$, and the maximal is for those elements laying on the line $x = 0$. Thus, $\overrightarrow{\delta}(A,B) = 15$. Now, for those elements of B satisfying $x \leq 42$, the distance is either 0 or $y - 25$. Hence, it is maximal for those elements laying on the line $y = 100$. For those satisfying $y \leq 25$, the distance is either 0 or $x - 42$. Hence, it is maximal for those elements laying on the line $x = 100$. For the remaining ones, i.e. those satisfying $x > 42$ and $y > 25$, the distance is given by $\sqrt{(x-42)^2 + (y-25)^2}$. Clearly, in this case, it is maximal for the point $(100,100)$ so that the overall forward distance is $\overrightarrow{\delta}(B,A) = \sqrt{(100-42)^2 + (100-25)^2} = \sqrt{8989} \approx 94.81$. Hence, we have that $\delta(A,B) \approx 94.81$.

Remark 5. Note that the considerations above also apply to the taxicab and the Chebyshev distance. More precisely, $\overrightarrow{\delta}(A,B) = 15$, while $\overrightarrow{\delta}(B,A)$ is still given by the distance between the point $(100,100)$ and $(42,25)$, which are $\mid 100 - 42 \mid + \mid 100 - 25 \mid = 133$ and $m + max\{\mid 100 - 42 - m \mid, \mid 100 - 25 - m \mid\} = 75$ for $m = min\{\mid 100 - 42 \mid, \mid 100 - 25 \mid\}$, respectively.

Example 9. Consider the opinions in the selection committee example in Section 4.2

$$B^0 = [\, c_s(5,6,3,5) \quad c_s(8,10,7,9) \quad c_s(5,7,2,4) \quad c'_s(1,4,6) \,]$$

Figure 2b shows the distances $\delta_{\geq s}(c_i,c_j)$ between agents i and j over time, along with the distance $\delta_{\geq s}(B^i)$ for the set of all the opinions when $s = 0.5$. Initially, Agents 2 and 3 exhibit the greatest divergence, whereas the opinions of Agents 1 and 3, as well as Agents 1 and 4, are more closely aligned. After the first interaction, the opinions of Agents 1 and 3 become more distant compared to the initial state, a trend also observed between Agents 3 and 4. Despite this, the overall distance $\delta_{\geq s}(B^i)$ decreases. After three iterations, all divergences disappear.

6   Concluding Remarks

We have introduced the Constraint Opinion Model, a generalization of the standard DeGroot model where opinions and influences are represented as soft constraints rather than single real values. Our framework allows for modeling belief revision scenarios involving partial information, uncertainty, and conditional influences. We illustrated the expressives of our approach through several examples and proposed a polarization measure for it based on a distance between two constraints.

Related Work. There is a great deal of work on generalizations and variants of the De-Groot model for more realistic scenarios (e.g., [1,2,12,9,10,3]). The work in [1] extends the DeGroot model to capture agents prone to confirmation bias, while [2] generalizes [1] by allowing agents to have arbitrary and differing cognitive biases. The study in [12] introduces a version of the DeGroot model in which self-influence changes over time, whereas influence on others remains constant. The works in [9,10] explore convergence and stability, respectively, in models where influences change over time. The study in [3] examines an asynchronous version of the DeGroot model. Nevertheless, to our knowledge, no generalizations of the DeGroot model address partial information.

This paper draws inspiration from the generalizations of constraint solving and programming [24] to deal with soft constraints representing preferences, probabilities, uncertainty, or fuzziness. We build on the *semiring*-based constraint framework [5,15], where (idempotent) semirings define the operations needed to *combine* soft constraints and determine when a constraint (or solution) is *better* than another. Other similar frameworks exist, such as the *valued* constraint framework [26], which has been shown to be equally expressive [7] to the one based on semirings.

Soft constraints [6] have been used to model agents or processes that share partial information in the style of concurrent constraint programming [25,21], where processes can *tell* (add) constraints to a common store of partial information, and synchronize by *asking* whether a constraint is entailed by the current store. Timed extensions of this framework were proposed later in [4], where agents can tell and ask different constraints along different time instants. In [23], these languages have been shown to have a strong connection with proof systems. To the best of our knowledge, this is the first time that (soft) constraints have been used in the context of belief revision.

Future Work. We have laid the foundations to study DeGroot-based belief revision under the lenses of soft constraints, where opinions and influences may include preferences, partial information and uncertainty. There are several directions to continue this work. The most obvious one is try to recover the standard theory of consensus in a generic semiring. This issue is not straightforward. Consider e.g. two elements X and Y such that $X \oplus Y = 1$ and $X \otimes Y = 0$, as they may be found in the free semiring given by the finite powerset. Now consider the matrix below

$$M = \begin{bmatrix} X & Y \\ Y & X \end{bmatrix}$$

The matrix is involutory, as e.g. M_3 in Example 2, hence it never stabilizes. However, differently from M_3, if both X and Y are different from 0, then the influence graph

is strongly connected, hence an immediate generalization of the consensus theorems does not hold. We believe that we could obtain it for semirings without zero divisors, along the line of the results on selective-invertible dioids presented in [17]. Second, it is worth exploring other models for social learning where, differently from the DeGroot model, not all the agents interact at the same time but only two (as in the gossip based model [14]) or some arbitrary set of them (as in the hybrid model in [20]). This is specially interesting when agents may discuss about different topics, as shown in Example 7. We also plan to extend the rewriting logic based framework proposed in [20] to the model proposed here, thus providing a framework for the (statistical) analysis of systems of agents discussing soft constraints. The present model allows for partial information about opinions and influence. Extending the DeGroot model with *partial information* about cognitive biases seem a natural line of future research.

References

1. Alvim, M.S., Amorim, B., Knight, S., Quintero, S., Valencia, F.: A formal model for polarization under confirmation bias in social networks. Logical Methods in Computer Science **19**(1) (2023). `https://doi.org/10.46298/LMCS-19(1:18)2023`
2. Alvim, M.S., da Silva, A.G., Knight, S., Valencia, F.: A multi-agent model for opinion evolution in social networks under cognitive biases. In: Castiglioni, V., Francalanza, A. (eds.) FORTE 2024. Lecture Notes in Computer Science, vol. 14678, pp. 3–19. Springer (2024). `https://doi.org/10.1007/978-3-031-62645-6_1`
3. Aranda, J., Betancourt, S., Díaz, J.F., Valencia, F.: Fairness and consensus in an asynchronous opinion model for social networks. In: Majumdar, R., Silva, A. (eds.) CONCUR 2024. LIPIcs, vol. 311, pp. 7:1–7:17. Schloss Dagstuhl - Leibniz-Zentrum für Informatik (2024). `https://doi.org/10.4230/LIPICS.CONCUR.2024.7`
4. Bistarelli, S., Gabbrielli, M., Meo, M.C., Santini, F.: Timed soft concurrent constraint programs. In: Lea, D., Zavattaro, G. (eds.) COORDINATION. Lecture Notes in Computer Science, vol. 5052, pp. 50–66. Springer (2008)
5. Bistarelli, S., Montanari, U., Rossi, F.: Semiring-based constraint satisfaction and optimization. Journal of ACM **44**(2), 201–236 (1997). `https://doi.org/10.1145/256303.256306`
6. Bistarelli, S., Montanari, U., Rossi, F.: Soft concurrent constraint programming. ACM Transactions on Computational Logic **7**(3), 563–589 (2006)
7. Bistarelli, S., Montanari, U., Rossi, F., Schiex, T., Verfaillie, G., Fargier, H.: Semiring-based csps and valued csps: Frameworks, properties, and comparison. Constraints **4**(3), 199–240 (1999)
8. Chandrasekhar, A.G., Larreguy, H., Xandri, J.P.: Testing models of social learning on networks: Evidence from a lab experiment in the field. Working Paper 21468, National Bureau of Economic Research (August 2015). `https://doi.org/10.3386/w21468`
9. Chatterjee, S., Seneta, E.: Towards consensus: Some convergence theorems on repeated averaging. Journal of Applied Probability **14**(1), 89–97 (1977), `https://doi.org/10.2307/3213262`
10. Chen, Z., Qin, J., Li, B., Qi, H., Buchhorn, P., Shi, G.: Dynamics of opinions with social biases. Automatica **106**, 374–383 (2019). `https://doi.org/https://doi.org/10.1016/j.automatica.2019.04.035`
11. DeGroot, M.H.: Reaching a consensus. Journal of the American Statistical association **69**(345), 118–121 (1974)

12. DeMarzo, P.M., et al.: Persuasion bias, social influence, and unidimensional opinions. The Quarterly Journal of Economics **118**(3), 909–968 (2003), `http://www.jstor.org/stable/25053927`

13. Esteban, J., Ray, D.: On the Measurement of Polarization. Econometrica **62**(4), 819–851 (July 1994), `https://ideas.repec.org/a/ecm/emetrp/v62y1994i4p819-51.html`

14. Fagnani, F., Zampieri, S.: Asymmetric randomized gossip algorithms for consensus. IFAC Proceedings Volumes (2008). `https://doi.org/10.3182/20080706-5-KR-1001.01528`

15. Gadducci, F., Santini, F., Pino, L.F., Valencia, F.D.: Observational and behavioural equivalences for soft concurrent constraint programming. Journal of Logic and Algebraic Methods in Programming **92**, 45–63 (2017). `https://doi.org/10.1016/J.JLAMP.2017.06.001`

16. Golub, B., Sadler, E.: Learning in social networks. SSRN (2017). `https://doi.org/10.2139/ssrn.2919146`

17. Gondran, M., Minoux, M.: Graphs, Dioids and Semirings. Springer (2008)

18. Hansson, S.O., Fermé, E.L., Cantwell, J., Falappa, M.A.: Credibility limited revision. Journal of Symbolic Logic **66**(4), 1581–1596 (2001). `https://doi.org/10.2307/2694963`

19. Makinson, D.: Screened revision. Theoria **63**(1-2), 14–23 (1997). `https://doi.org/10.1111/j.1755-2567.1997.tb00737.x`

20. Olarte, C., Ramírez, C., Rocha, C., Valencia, F.: Unified opinion dynamic modeling as concurrent set relations in rewriting logic. In: Ogata, K., Martí-Oliet, N. (eds.) WRLA 2024. Lecture Notes in Computer Science, vol. 14953, pp. 104–123. Springer (2024). `https://doi.org/10.1007/978-3-031-65941-6_6`

21. Olarte, C., Rueda, C., Valencia, F.D.: Models and emerging trends of concurrent constraint programming. Constraints **18**(4), 535–578 (2013). `https://doi.org/10.1007/S10601-013-9145-3`

22. Paz, J., Rocha, C., Tobòn, L., Valencia, F.: Consensus in Models for Opinion Dynamics with Generalized-Bias. In: 13th International Conference on Complex Networks and Their Applications. Lecture Notes in Computer Science, vol. 12062, pp. 253–269. Springer, Istanbul, Turkey (Dec 2024). `https://doi.org/10.48550/arXiv.2409.10809`, `https://inria.hal.science/hal-04918975`

23. Pimentel, E., Olarte, C., Nigam, V.: A proof theoretic study of soft concurrent constraint programming. Theory and Practice of Logic Programming **14**(4-5), 649–663 (2014). `https://doi.org/10.1017/S147106841400026X`

24. Rossi, F., van Beek, P., Walsh, T. (eds.): Handbook of Constraint Programming, Foundations of Artificial Intelligence, vol. 2. Elsevier (2006)

25. Saraswat, V.A.: Concurrent Constraint Programming. MIT Press (1993)

26. Schiex, T., Fargier, H., Verfaillie, G.: Valued constraint satisfaction problems: Hard and easy problems. In: IJCAI (1). pp. 631–639 (1995)

27. Xia, W., Ye, M., Liu, J., Cao, M., Sun, X.M.: Analysis of a nonlinear opinion dynamics model with biased assimilation. Automatica **120**, 109113 (2020). `https://doi.org/https://doi.org/10.1016/j.automatica.2020.109113`